

Bhavika Baddur

baddurbhavika@gmail.com — +971558365400 — [LinkedIn](#) — [GitHub](#)

Nationality: Indian — Visa Status: Residency Visa

Professional Summary

Computer Science graduate (GPA: 8.9/10) with hands-on experience in AI systems engineering and data engineering at Wissen Infotech. Built production-grade multi-agent AI platforms using LangGraph, MCP, Supabase, and Power BI, alongside Microsoft Azure data pipelines, Microsoft Fabric solutions, and Copilot Studio agents. Completed end-to-end projects spanning agentic AI orchestration, data pipelines, and BI dashboards, demonstrating strong problem-solving, cloud, and full-stack data development skills. Eager to contribute to AI and data-driven solutions in Dubai's technology and transport sector.

Skills Summary

- **Languages:** Python, Java, C, JavaScript, TypeScript, SQL
- **AI / ML & Agentic Systems:** LangGraph, LangChain, MCP (Model Context Protocol), RAG, FAISS, Llama 3, Prompt Engineering, NLP
- **Data Engineering:** PySpark, Databricks, Delta Lake, Microsoft SQL Server, Microsoft Fabric, Power BI, Pandas
- **Cloud & Automation:** Microsoft Azure, Copilot Studio, Power Automate, n8n
- **Tools & Platforms:** Git, GitHub, Docker, Streamlit

Work Experience

AI & Data Engineering

Feb 2026 – Present

Wissen Infotech, Hyderabad

- Created and launched an AI-powered Telugu FIR translation platform using **Tesseract OCR** and **Sarvam AI APIs**, building an end-to-end Streamlit app that simplified the document translation process.
- Designed a dual-source analytics chatbot with **Copilot Studio** that mapped custom SQL logic over Azure SQL and Excel datasets (26,000+ transactions), making it easy for business users to query data without writing code.
- Automated data quality checks by building **Power Automate** workflows across Azure SQL and Excel files to automatically handle duplicate filtering, **NULL** analysis, and **Pass/Fail** report generation.
- Set up a high-performance Microsoft Fabric data ingestion pipeline between ADLS accounts and configured Fabric Mirroring for continuous, near-real-time Azure SQL database synchronisation.

Projects

1. FinSight Enterprise AI - Multi-Agent Financial Intelligence Platform — [Link](#) — [Live Demo](#)

Stack: LangGraph, MCP, n8n, Supabase, Qdrant, Groq (Llama 3.3 70B), Power BI, PySpark, Streamlit, Tesseract OCR

- Built an end-to-end financial document intelligence system processing 387+ transactions across multiple business entities, combining a production-grade **Medallion data pipeline** (Bronze → Silver → Gold) with a 5-agent **LangGraph** orchestration layer for automated risk analysis and anomaly detection.
- Engineered 4 custom **MCP (Model Context Protocol) tool servers** mediating all database interactions, enforcing a strict separation between AI agents and data access for governance and auditability.
- Developed an **Intelligent Document Processing (IDP)** pipeline using Tesseract OCR and LLM-based extraction with Pydantic validation, achieving 88-100% field-level confidence on bank statements and invoices.
- Designed interactive **Power BI dashboards** connected directly to the Gold layer, alongside a natural-language-to-SQL agent enabling business users to query transaction data without writing code.
- Implemented **human-in-the-loop governance** with automated risk-based interrupts, an immutable audit trail, and **n8n** workflow automation, achieving full observability via LangSmith tracing and a custom LLMops cost-tracking dashboard.

2. DataForge AI - Multi-Agent Analytics Copilot — [Link](#)

Stack: LangGraph, LangChain, Groq (Llama 3.3 70B), FAISS, Microsoft SQL Server, Streamlit, Plotly, PySpark

- Developed an agentic intelligence layer leveraging LangGraph to dynamically translate complex natural-language business requests into executable **SQL queries, business insights, and PySpark ETL transformations**.
- Architected a 7-agent orchestration framework featuring a specialised **Context Agent** for multi-turn conversation memory and an intent-based **Planner Agent** to route workflows with minimal latency.
- Deployed low-latency inference using Groq (Llama 3.3 70B) integrated with a **FAISS-powered RAG system** to enforce strict database schema-awareness and inject safety guardrails against unsupported business logic.
- Established an automated **LLMops evaluation pipeline** to systematically benchmark and validate planner routing accuracy, vector store retrieval coverage, and synthetic code-generation safety.

3. RetailIQ - End-to-End Retail Intelligence Lakehouse — [Link](#)

Stack: Python, PySpark, Databricks, Delta Lake, Microsoft SQL Server, Power BI, Pandas

- Engineered a production-ready Medallion Lakehouse architecture (**Bronze** → **Silver** → **Gold**) utilising Python and PySpark to ingest and process transactional retail data across 4 regions.
- Designed and deployed a **Star-Schema Data Warehouse** in MS SQL Server, writing modular Python data loaders to automate dimension and fact table formatting.
- Optimised the pipeline natively within **Databricks using PySpark and Delta Lake**, implementing centralised logging and type-casting to deliver 5 optimised Gold-layer KPI aggregation tables.
- Developed an interactive Power BI dashboard mapped to the analytical warehouse layers, transforming raw data into automated visual insights for regional revenue and product trends.

4. Hinglish Sentiment Analysis - Multilingual E-commerce NLP — [Link](#)

- Developed a code-mixed NLP pipeline for Hindi-English product reviews, tackling transliteration noise, informal grammar, and mixed-script text that standard models cannot handle.
- Built a custom preprocessing pipeline covering normalisation, denoising, and tokenisation; compared traditional ML models with **XLM-ROBERTa** on the same dataset.
- Redesigned classification from 3-class to **binary sentiment**, improving accuracy to **77%** and reducing label ambiguity; deployed a Streamlit dashboard for real-time review analysis.
- Surfaced aspect-level customer insights across product quality, delivery, and pricing, enabling product teams to identify actionable improvement areas from unstructured review data.

5. Resume Matching System - NLP-Based HR Tech — [Link](#)

- Built an **NLP-powered** resume-to-job-description matching engine, evaluated on 300+ resumes, improving candidate-shortlisting precision by **15%** over keyword-only baseline approaches.
- Applied **TF-IDF** and semantic similarity techniques to rank candidates; designed for scalability to handle high-volume applicant pipelines common in enterprise hiring workflows.

Education

KG Reddy College of Engineering and Technology, Hyderabad
B.Tech, Computer Science and Engineering

August 2026
GPA: 8.9/10

Certifications

- Data Analysis with Python – IBM SkillsBuild
- Introduction to Artificial Intelligence – Infosys Springboard