

Bhavika Baddur

baddurbhavika@gmail.com • +91 807443133 • [LinkedIn](#) • [GitHub](#)

EDUCATION

KG Reddy College of Engineering and Technology, Hyderabad
B.Tech, Computer Science and Engineering

August 2026
GPA: 8.9/10

TECHNICAL SKILLS

- **Languages:** Python, Java, C, JavaScript, TypeScript, SQL
- **AI / ML & Agentic Systems:** LangGraph, LangChain, MCP (Model Context Protocol), RAG, FAISS, Llama 3, Prompt Engineering, NLP
- **Data Engineering:** PySpark, Databricks, Delta Lake, Microsoft SQL Server, Microsoft Fabric, Power BI, Pandas
- **Cloud & Automation:** Microsoft Azure, Copilot Studio, Power Automate, n8n
- **Tools & Platforms:** Git, GitHub, Docker, Streamlit

WORK EXPERIENCE

AI & Data Engineering Intern
Wissen Infotech, Hyderabad

Feb 2026 – May 2026

- Built and deployed an AI-powered Telugu FIR translation platform using **Tesseract OCR** and **Sarvam AI API**, reducing manual document translation effort through an end-to-end **Streamlit workflow**.
- Developed a natural language analytics chatbot over **Azure SQL** and **Excel datasets** containing 26,000+ transactions, enabling business users to retrieve insights without writing SQL using **Copilot Studio** and **Power Automate**.
- Architected a **Microsoft Fabric** ADLS-to-ADLS ingestion pipeline with **Fabric Mirroring** for near-real-time **Azure SQL replication** and continuous data synchronisation.

PROJECTS

1. **FinSight Enterprise AI - Multi-Agent Financial Intelligence Platform** [Link](#)
Stack: LangGraph, MCP, n8n, Supabase, Qdrant, Groq (Llama 3.3 70B), PySpark, Streamlit, Tesseract OCR

- Built an end-to-end financial document intelligence system processing 387+ transactions, combining a **Medallion data pipeline** (Bronze → Silver → Gold) with a 5-agent **LangGraph** orchestration layer for automated risk analysis and anomaly detection.
- Engineered 4 custom **MCP (Model Context Protocol) tool servers** for governed database access and an **IDP pipeline** (Tesseract OCR + LLM extraction) achieving 88-100% field-level confidence.
- Implemented **human-in-the-loop governance** with risk-based interrupts and **n8n** workflow automation, achieving full observability via LangSmith tracing.

2. **DataForge AI - Multi-Agent Analytics Copilot** [Link](#)
Stack: LangGraph, LangChain, Groq (Llama 3.3 70B), FAISS, Microsoft SQL Server, Streamlit, Plotly, PySpark

- Developed an agentic intelligence layer leveraging **LangGraph** to dynamically translate complex natural-language business requests into executable **SQL queries**, business insights, and **PySpark ETL** transformations.
- Architected a 7-agent orchestration framework featuring a specialised **Context Agent** for multi-turn conversation memory and an intent-based **Planner Agent** to route workflows with minimal latency.
- Deployed low-latency inference using **Groq (Llama 3.3 70B)** integrated with a **FAISS-powered RAG** system to enforce strict database schema-awareness and inject safety guardrails against unsupported business logic.
- Established an automated **LLMOps evaluation pipeline** to systematically benchmark and validate planner routing accuracy, vector store retrieval coverage, and synthetic code-generation safety.

3. **RetailIQ - End-to-End Retail Intelligence Lakehouse** [Link](#)
Stack: Python, PySpark, Databricks, Delta Lake, Microsoft SQL Server, Power BI, Pandas

- Engineered a production-ready **Medallion Lakehouse architecture** (Bronze → Silver → Gold) utilising **Python** and **PySpark** to ingest and process transactional retail data across 4 regions.
- Designed and deployed a **Star-Schema Data Warehouse** in **MS SQL Server**, writing modular Python data loaders to automate dimension and fact table formatting.
- Optimised the pipeline natively within **Databricks** using **PySpark** and **Delta Lake**, implementing centralised logging and type-casting to deliver 5 optimised Gold-layer **KPI aggregation tables**.
- Developed an interactive **Power BI dashboard** mapped to the analytical warehouse layers, transforming raw data into automated visual insights for regional revenue and product trends.

4. **Hinglish Sentiment Analysis - Multilingual E-commerce NLP** [Link](#)

- Developed a code-mixed **NLP pipeline** for Hindi-English product reviews handling transliteration noise, informal grammar, and mixed-script text.
- Compared traditional **ML models** with **XLM-ROBERTa**, achieving 77% sentiment classification accuracy on code-mixed review datasets.
- Built a **Streamlit dashboard** surfacing aspect-level insights across product quality, delivery, and pricing for real-time review analysis.

CERTIFICATIONS

- **Data Analysis with Python** – IBM Skills Build • **Introduction to Artificial Intelligence** – Infosys Springboard